

6. Transformer Tokenizers

Generative Algorithms for Sound and Music



Universitat
Pompeu Fabra
Barcelona

MTG

Music Technology
Group

Why I am teaching this class

- Spent a whole year dissecting and reimplementing MIDITok
- I <3 u guys

What is a tokenizer?

- Translation layer between the transformer and its modality
- Encodes
- Decodes
- Is invertible*

Some terms

- Vocabulary
 - List of all recognized tokens
- Encode / Decode
 - Move from medium to tokens, and back.
- Token
 - Discrete element of information
- Id
 - Associated integer to a token

Why do we need a tokenizer anyway?

- Transformers operate on discrete data
- We need a reliable, reversible way to encode
- We don't have infinite context lengths or resources
 - Need to be efficient and concise



Tokenizers in NLP (Text-based models)

- The “foundation” is the set of all characters
- There are different philosophies and approaches to this problem

Tokenizer Philosophies in NLP

- Dictionary-based
- Character-based
- Byte-based
- Rule-based: Split dataset given pre-defined rules.
- Statistics-based: Start with a set of base tokens, create groupings based on recurrent patterns

Statistics-based tokenization in NLP

Start with a base vocabulary, add new tokens based on recurrent patterns

>>> What is a good starter set?

- Byte Pair Encoding
- WordPiece
- Any other methods?



Text Tokenizers

Music Tokenizers

What musical information do we want to tokenize?

Initial Considerations

- Balance between available data, token variety, vocabulary size
- Purpose of model: What are we generating?


Nooo, you have
to know how MIDI
works



It doesn't matter
how MIDI works



It doesn't matter
how MIDI works



A History of MIDI

- Started as a communication between Synthesisers (1983)
- Need for saving files came later!
- MIDI file standard is more of a guideline...

MIDI files - Header

- Defines NumTracks, Division
- Division = Ticks per quarter note
 - Typical values: 384, 480, 960, 1920... (Why?)

MIDI files - Events

- Note Event: On / Off, Channel, Pitch, Velocity
- CC messages: Channel, ID, Value
- Ticks: Time in ticks between next event and previous
- Meta Events: TimeSig, Tempo, KeySig, Or even Text
- Track Events
- SysEx Events

MIDI files - What's not in it?

- Explicit bar numbers
- Beat emphasis
- Phrasing information

All meta events are optional - Cannot assume they are always present.

What you need to know about MIDI files

- Two types: Performance, and Score based
- Notes are encoded with NoteOn and NoteOff messages
- Events are separated by Ticks
- Can include CC messages and even text

MIDI files - Implications for Tokenizers

- Don't have necessarily the information we want
- Need to consider the quality of the files, not just quantity
- Tick-based timing needs processing
- May have information we do not want!

Tokenizing notes

- Crucially, Notes have a pitch
- Definite Pitch vs Pitch Intervals?
- Duration vs TimeShift + NoteOff?
- Velocity Token?
- Chord Token?


Tokenizing time

- Duration
- Rest
- TimeShift
- Position
- Bar
- MicroTiming?


Tokenizing Meta Events

- Tempo
- Time Signatures
- PitchBends
- Pedals
- Program
- Track
- CC?

MIDI-like



Revamped MIDI



Tokenizing Polyphony


Hierarchical - Program token
before each note

- Longer token sequence
- Notes encoded sporadically
- Simultaneous notes are close together
- Interrupts melodic flow

Flat - Track Separators every 1-2 bars

- Shorter token Sequence
- Organized notes
- Simultaneous notes are far apart
- Interrupts harmonic flow

Revamped MIDI Plus



To Pool or Not To Pool?

Pooling = Process of merging tokens into new tokens

- Can be done statistically: BPE, WordPiece...
 - WordPiece equivalent for music?
- Can be done algorithmically:
 - Can group all note-related tokens in a pool

What goes in a pool? Is it necessary? Can it hurt?

Octuple

Time Sig 4/4	Time Sig 4/4	Time Sig 4/4	Time Sig 4/4
Tempo 116	Tempo 116	Tempo 116	Tempo 116
Bar 0	Bar 0	Bar 0	Bar 1
Position 0	Position 7	Position 27	Position 0
Program 0	Program 0	Program 0	Program 0
Duration 1.0	Duration 1.0	Duration 0.4	Duration 4.0
Velocity 80	Velocity 84	Velocity 72	Velocity 80
Pitch D3	Pitch A3	Pitch G3	Pitch A3



MuMIDI

			Duration 1.0		Duration 1.0		Duration 0.4		Duration 4.0	
			Velocity 80		Velocity 84		Velocity 72		Velocity 80	
Tempo 116	Tempo 116	Tempo 116	Tempo 116	Tempo 116	Tempo 116	Tempo 116	Tempo 116	Tempo 116	Tempo 116	
Pos None	Pos 0	Pos 0	Pos 0	Pos 7	Pos 7	Pos 27	Pos 27	Pos None	Pos 0	Pos 0
Bar 0	Bar 0	Bar 0	Bar 0	Bar 0	Bar 0	Bar 0	Bar 0	Bar 1	Bar 0	Bar 0
Bar	Position 0	Program	Pitch D3	Position 7	Pitch A3	Position 27	Pitch G3	Bar	Position 0	Pitch A3



Special Tokens

Special Tokens = Indirect information, don't have a concrete symbol.

- You know some already: [BOS, EOS, PAD, MASK]
- Steer the direction of the generations
- When generated, we ignore them
- Any ideas?

Adjusting for misplaced generated tokens

- Transformers are unpredictable
- Need to have predefined rules for failure cases
- Some failures are worse than others:
 - Missing Velocity token?
 - Missing Pitch token?
 - Early / Late / Missing Bar token?
- If you are getting too many, your token resolution may not be adequate to your data

Activity: Explore phrase-based pooling

If Notes are to music what letters are to speech, it would stand to reason that musical phrase-based pooling could improve our tokenizers, as seen in NLP.

- What's a phrase?
- Can we tokenize all phrases?
- Notes are already quite complex; what info do we want in the pool, and what not?
- Rule-based pooling or statistics-based?
- Pick a fragment of sheet music and tokenize it by hand

In clusters of 4 - 6 people

MIDITok

- Python Library - Compilation of tokenizer implementations
- Built on top of SyMusic for Score-based representations of MIDI
- HuggingFace hub integration
- Makes it easy to write your own tokenizer within the framework

mDITok Implemented Tokenizers

- [link](#)

Why implement your tokenizer in **MIDITok**

- Handles Saving - Loading of your configs
- Provides helpful methods to use
- Easy training for BPE
- Clean, with pre-built pipelines for testing
- Built-in token validation
- Can act as a robust boilerplate